


# Scenario For Web Development



<https://brtneura.com>

## Copyright Information

This document is the exclusive property of BRTNeura Technology LLP; the recipient agrees that they may not copy, transmit, use, or disclose the confidential and proprietary information in this document by any means without the expressed and written consent of BRTNeura. By accepting a copy, the recipient agrees to adhere to these conditions to the confidentiality of BRTNeura's practice and procedures.

## Document Control

Term	Details
Document Title	Candidates Exam Scenario
Author	BRTNeura
Version Number	
Date Created	
Date Submitted	

## Index

1. Executive Summary.....	4
2. Scenario.....	5
3. Evaluation Criteria.....	5
4. Tips to Pass (and score 85+).....	8
5. Suggested 6–8h game plan .....	9
6. Common pitfalls (avoid!) .....	9

# 1. Executive Summary

## What you'll do:

You'll build a small **full-stack web app** based on **one scenario** (assigned to you). It must include:

- **Backend:** Python **FastAPI** (Pydantic, Uvicorn)
- **Frontend:** **React + TypeScript** (Vite)
- **Auth:** **Firebase Auth** (server verifies tokens)
- **DB:** As specified in your scenario (**Firestore / Postgres / SQLite**)
- **AI/Analytics:** A small, **simple** component (e.g., heuristic or basic model) with a quick evaluation
- **Tests:**  $\geq 3$  **backend** (pytest) +  $\geq 1$  **frontend** (React Testing Library)
- **Docs & DevEx:** README with setup, **OpenAPI** at /docs (or Postman), architecture diagram, .env.example, and a  **$\leq 5$ -min Loom** walkthrough

**Timebox:** 6–8 hours.

**Pass bar:** Score  $\geq 70/100$  on the common rubric. **Top tier:** 85+

## What to submit (deliverables checklist)

- Repo with two folders: /api (FastAPI) and /web (React)
- README.md with:
  - Setup steps & run commands (dev, test, seed)
  - **Architecture diagram** (Mermaid or PNG)
  - Endpoints + sample requests/responses
  - **ADR** (why you chose your datastore/indices)
  - **"How I used Cursor"** (3–5 lines)
- **OpenAPI** at /docs (or a Postman/Thunder collection)
- **Seed script** to load sample data ( $\geq 1k$  rows where relevant)
- **Tests** (all pass with one command)
- **Deploy links** (Netlify/Vercel for FE + Render/Fly/Cloud Run for API) **or** Docker Compose
- **Loom video** ( $\leq 5$  min) demoing the core flow

## 2. Scenario

### Task Runner & Webhook Orchestrator (Postgres)

You are a developer who needs to develop this website whose function is to let users define HTTP webhooks to run on schedules and view run logs with retries and a dead-letter queue.

**Must-use tech:**

FastAPI, React+TS, Firebase Auth, Postgres, PyTest, background worker.

**Schema:**

tasks(id, user\_id, url, method, headers\_encrypted, body, schedule\_cron, enabled)

runs(id, task\_id, status, latency\_ms, response\_code, error, created\_at)

dlq(id, task\_id, error, created\_at)

**API:**

- POST /tasks GET /tasks GET /tasks/{id}/runs
- Worker executes tasks, retries with exponential backoff up to 3, then DLQ.

**Frontend:**

Task composer, run history table with filters, latency histogram.

**AI:**

Failure classifier (network/auth/payload) by simple rules + explanation string.

**Security:**

Encrypt headers server-side; mask in UI.

**Tests:**

(1) Retry/backoff sequence, (2) DLQ path, (3) Secret masking not leaked.

## 3. Evaluation Criteria

**TOTAL Score: 100**

### 1) Functional completeness (15)

- **Core user flows work** end-to-end as specified (create/read/update, scenario-specific flows). (0–8)
- **Edge cases handled** (empty state, invalid input, not found). (0–4)
- **Polish** (pagination, meaningful errors/toasts). (0–3)

### 2) Architecture & data modeling (10)

- **Sound schema/collections** with correct keys, indexes, relations. (0–5)
- **Separation of concerns** (routers/services/models/utils). (0–3)
- **Trade-offs explained** in ADR (why SQL vs Firestore, why chosen indexes). (0–2)

### 3) Backend API quality (15)

- **Contract** matches scenario (routes, verbs, status codes, validation). (0–6)
- **Performance awareness** (pagination server-side, N+1 avoided, basic indexing). (0–5)
- **Resilience** (timeouts, retries/background jobs where required). (0–4)

### 4) Frontend implementation (10)

- **State management** is reasonable (React Query/contexts/hooks). (0–4)
- **Forms & UX** (validation, loading/error states, accessible components). (0–4)
- **Integration** with backend/auth done correctly. (0–2)

### 5) Auth & security (10)

- **Firebase Auth implemented**; tokens verified **server-side**. (0–4)
- **Authorization** (role checks, multi-tenant scoping where applicable). (0–4)
- **Secrets & inputs** (no secrets in client, input sanitization, rate limits if required). (0–2)

### 6) AI/analytics piece (10)

- **Exists and is relevant** to the scenario (not a stub). (0–4)
- **Basic evaluation** (e.g., confusion matrix, precision@k, z-score thresholds, or documented assumptions). (0–4)
- **Cost/complexity discipline** (simple, explainable approach; no paid APIs). (0–2)

### 7) Performance & reliability (10)

- **Seeded data size** ( $\geq 1k$  rows/items where relevant). (0–3)
- **Latency** ( $P50 \leq 300$  ms on core list endpoints locally—documented). (0–3)

- **Background work** (schedulers/workers behave as specified; retries/backoff). (0–4)

#### 8) Testing (10)

- **Backend tests ( $\geq 3$ )** cover a happy path and at least one failure path. (0–6)
- **Frontend test ( $\geq 1$ )** covers a key interaction (filter, submit, or route). (0–3)
- **Tests pass** via a single command. (0–1)

#### 9) Code quality & DevEx (5)

- **Readable code** (names, small functions, lint/format). (0–3)
- **DX scripts** (dev, test, seed, or Makefile). (0–2)

#### 10) Documentation & communication (5)

- **README:** setup, run, env, architecture diagram, endpoints, “How I used Cursor” note. (0–4)
- **Postman/OpenAPI** or collection linked; short Loom  $\leq 5$  min. (0–1)

### Bonus & penalties ( $\pm 10$ )

#### Bonus (+1–10):

- Stretch feature shipped well (+2–4)
- Observability (health, structured logs, basic tracing) (+1–2)
- Clean migrations/data seeds for SQL; Firestore rules written (+1–2)
- Thoughtful UX touches (keyboard nav, a11y) (+1–2)

#### Penalties (–1–10):

- Hard-coded secrets or client-side secrets (–5)
- Auth missing on server (–5)
- Can’t run from README (–3)
- Heavy copy-paste/boilerplate without integration (–2)
- Fails scenario contract (wrong routes/DB) (–3)

#### Auto-reject conditions

- No server-side auth/authorization when required.
- App cannot be started following README.
- Plagiarism or inability to explain code during follow-up.

- Data exposure (e.g., cross-tenant leaks).

## 4. Tips to Pass (and score 85+)

### 1) Land the core user flow early

- Implement **one happy path** E2E in the first 2–3 hours (auth → create → list → detail).
- Defer fancy UI until after the flow works.

### 2) Keep AI simple but *real*

- Use **rules + a tiny model** (e.g., logistic regression, TF-IDF + cosine).
- Show **one tiny metric** (confusion matrix, precision@k, or threshold rationale).
- No paid APIs needed.

### 3) Prove server-side auth

- Validate Firebase ID token **in FastAPI** (every protected route).
- If multi-tenant: enforce `org_id` on the server for every query.

### 4) Nail pagination & indexing

- **Server-side pagination** (limit + cursor/offset) is required on “list” endpoints.
- Add 1–2 **indexes** that actually speed up your list queries; note them in the README.

### 5) Ship tests that matter

- Backend: (a) happy path, (b) **one failure path** (e.g., validation/rate-limit), (c) a behavior tied to your scenario (e.g., dedupe, scheduler).
- Frontend: test a **key interaction** (filtering, submit, or rendering a list).

### 6) Be reliable under load (seed data)

- Seed **≥1k rows** (transactions, leads, products... as relevant).
- Paste **measured p50 latency** (curl or simple timer) into README.

### 7) Keep secrets secret

- Provide **.env.example**; never hardcode keys; never put secrets in the client bundle.

### 8) Background jobs: make them observable

- Add a **health endpoint** and log each run with status/latency (for schedulers/workers).
- Document retry/backoff behavior if your scenario calls for it.



## 9) Communicate like a pro

- Short, accurate **README**; clear limitations; “what I’d do next.”
- Loom shows: auth → core flow → the AI/analytics bit → tests passing (10–15 sec each).

## 10) Use Cursor wisely (and say how)

- Show *where* you used it (e.g., stub generation, test scaffolding) and *how* you verified output.
- Don’t accept code blindly—commit in **small diffs** with messages.

## 5. Suggested 6–8h game plan

1. **Setup (45 min)**: Repos, auth stub, DB schema/collections, routes/components.
2. **Core flow (2–3h)**: E2E path working + server-side pagination.
3. **AI/analytics (60–90 min)**: Implement simplest viable logic + one metric.
4. **Polish + perf (60 min)**: Seed ≥1k rows; add indexes; measure p50; fix logs.
5. **Tests (45–60 min)**: 3 backend + 1 frontend, all green.
6. **Docs & demo (30–45 min)**: README, OpenAPI, Loom.

## 6. Common pitfalls (avoid!)

- Client-only auth (no server verification)
- No pagination; “load all” queries
- Hardcoded secrets; missing .env.example
- Can’t run from README; missing scripts
- AI part is a placeholder (no metric, no effect on UI)
- No seed data → you can’t show performance